

Swarm: A Lightweight yet Highly Effective Method for Improving Random Testing

Alex Groce, *Member, IEEE*, Chaoqiang Zhang, Yang Chen, Eric Eide, Mohammad Amin Alipour, and John Regehr *Member, IEEE*,



Abstract—Swarm testing is an approach to improving the effectiveness of random testing that replaces the traditional testing procedure of generating tests from a single probability distribution with the use of a “swarm” of distributions each of which disallows certain test behaviors completely. This paper shows that swarm has two highly desirable properties. First, swarm is in many cases extremely effective, resulting in 40% or better improvements in *distinct faults detected* for critical real-world systems such as compilers. Second, swarm is an encouragingly *lightweight* method for improving random testing, with very few demands on would-be users; applying swarm to an existing random testing system is almost always trivially easy. In our experience, even a programmer unfamiliar with a complex testing system can often apply swarm testing to it. The widespread applicability and effectiveness of swarm testing suggests a new domain of random testing research: the search for techniques that rely only on general features of most test spaces, and do not require much additional programmer effort or analytical machinery to apply to already existing random testing frameworks, which are frequently complex and difficult to modify. This paper presents case studies demonstrating the lightweight nature and high effectiveness of swarm across a set of important real-world case studies.

This paper focuses on answering a single question: In random testing, can a diverse set of *testing configurations* perform better than a single, possibly “optimal” configuration? An example of a test configuration would be, for example, a list of API calls that can be included in test cases. Conventional wisdom in random testing [1] has assumed a policy of finding a “good” configuration and running as many tests as possible with that configuration. Considerable research effort has been devoted to the question of how to tune a “good configuration,” e.g., how to use genetic algorithms to optimize the *frequency* of various method calls [2], or how to choose a length for tests [3]. As a rule, the notion that some test configurations are “good” and that finding a good (if not truly optimal, given the size of the search space) configuration is important has not been challenged. Furthermore, in the interests of maximizing coverage and fault detection, it has been assumed that

a good random test configuration includes as many API calls or other input domain features as possible, and this has been the guiding principle in large-scale efforts to test C compilers [4], file systems [5], and utility libraries [6]. The rare exceptions to this rule have been cases where a feature makes tests too difficult to evaluate or slow to execute, or when static analysis or hand inspection can demonstrate that an API call is unrelated to state [5]. For example, including pointer assertions may make compiling random C programs too slow with some compilers.

In general, if a call or feature is omitted from some tests, it is usually omitted from all tests. This approach seems to make intuitive sense: omitting features, unless it is necessary, means *giving up on detecting some faults*. However, this objection to feature omission only holds so long as testing is performed using a single test configuration. Swarm testing, in contrast, uses a diverse “swarm” of test configurations, each of which *deliberately omits certain API calls or input features*. As a result, given a fixed testing budget, swarm testing tends to test a more diverse set of inputs than would be tested under a so-called “optimal” configuration (perhaps better referred to as a *default* configuration) in which every feature is available for use by every test.

One can visualize the impact of swarm testing by imagining a “test space” defined by the contents of tests. As a simple example, consider testing an implementation of a stack ADT that provides two operations, push and pop. One can visualize the test space for the stack ADT using these features as axes: each test is characterized by the number of times it invokes each operation. Any method for randomly generating test cases results in a probability distribution over the test space, with the value at each point (x, y) giving the probability that a given test will contain exactly x pushes and y pops (in any order). To make this example more interesting, imagine the stack implementation has a capacity bug, and will crash whenever the stack is required to hold more than 32 items.

Missing fig illustrates the situation for testing the stack with a test generator that chooses pushes and pops with equal probability. The generator randomly chooses an

- A. Groce, C. Zhang, and M. Alipour are with the School of Electrical Engineering and Computer Science, Oregon State University, Corvallis, OR, 97331.
- Y. Chen, E. Eide, and J. Regehr are with the School of Computing, University of Utah, Salt Lake City, 84112.

input length and then decides if each operation is a push or a pop. The graph shows the distribution of tests produced by this generator over the test space. The graph also shows contour lines for significant regions of the test space. Where $P_{fail} = 1$, a test chosen randomly from that region is certain to trigger the stack’s capacity bug; where $P_{fail} = 0$, no test can trigger the bug. As missing fig shows, this generator only rarely produces test cases that can trigger the bug.

Now consider a test generator based on swarm testing. This generator first chooses a non-empty subset of the stack API and then generates a test case using that subset. Thus, one-third of the test cases contain both pushes and pops, one-third just pushes, and one-third just pops. ?? shows the distribution of test cases output by this generator. As is evident from the graph, this generator often produces test cases that trigger the capacity bug.

Although simple, this example illustrates the dynamics that make swarm testing work. The low dimensionality of the stack example is contrived, of course, and we certainly believe that programmers should make explicit efforts to test boundary conditions. As evidenced by the results presented in this paper, however, swarm testing generalizes to real situations in which there may be dozens of features that can be independently turned on or off. It also generalizes to testing real software in which faults are very well hidden.

Every test generated by any swarm configuration can, in principle, be generated by a test configuration with all features enabled. However—as the stack example illustrates—the probability of covering parts of the state space and detecting certain faults can be demonstrably higher when a diverse set of configurations is tested.

Swarm testing has several important advantages. First, it is low cost: in our experience, existing random test case generators already support or can be easily adapted to support feature omission. Second, swarm testing reduces the amount of human effort that must be devoted to tuning the random tester. In our experience, tuning is a significant ongoing burden. Finally—and most importantly—swarm testing makes significantly better use of a fixed CPU time budget than does random testing using a single test configuration, in terms of both coverage and fault detection. For example, we performed an experiment where two machines, differing only in that one used swarm testing and one did not, used Csmith [4] to generate tests for a collection of production-quality C compiler versions for x86-64. During one week of testing, the swarm machine found 104 distinct ways to crash compilers in the test suite whereas the other machine—running the default Csmith test configuration, which enables all features—found only 73. An improvement of more than 40% in terms of number of bugs found, using a random tester that has been intensively tuned for several years, is surprising and significant.

Even more surprising were some of the details. We found, for example, a compiler bug that could only be triggered by programs containing pointers, but which

was almost never triggered by inputs that contained arrays. This is odd because pointer dereferences and array accesses are very nearly the same thing in C.¹ Moreover, we found another bug in the same compiler that was only triggered by programs containing arrays, but which was almost never triggered by inputs containing pointers. Fundamentally, it appears that omitting features while generating random test cases can lead to improved test effectiveness.

Our contributions are as follows. First, we characterize *swarm testing*, a pragmatic variant of random testing that increases the diversity of generated test cases with little implementation effort. The swarm approach to diversity differs from previous methods in that it focuses solely on *feature omission diversity*: variance in which possible input features are *not* present in test cases. Second, we show that—in three case studies—swarm testing offers improved coverage and bug-finding power. Third, we offer some explanations as to *why* swarm testing works.

1 SPIDERMONKEY RESULTS

Figures 1-3 show a comparison between swarm testing and a default strategy for Mozilla’s SpiderMonkey JavaScript engine, for release versions 1.6, 1.7, and 1.8.5. Tests were generated using the last public version of the jsfunfuzz JavaScript fuzzing tool [7], both in its original form and modified to perform swarm testing. The modification of jsfunfuzz to perform swarm testing was extremely simple. First, the source code was reformatted with a regular expression to place individual choices in the random selectors on separate lines. Second, the reformatted code was “swarmed” by a 30 line python script that removed items in each random choice, with 50% probability for each removal. We estimate that this process took perhaps 30 minutes for a developer unfamiliar with either the jsfunfuzz code or the JavaScript language.

1.1 Experimental Procedure

Tests were generated in 30 minute budget runs, and coverage and faults or failures were computed for each 30 minute run. The graphs show cumulative branch coverage and faults or failures for 48 hours of testing. In addition to a pure-swarm and pure-default strategy, we also show results for a strategy that performs swarm testing for 24 hours, then switches to the default configuration, in the hope that this will explore behaviors hard to reach under swarm.

Distinct faults for versions 1.6 and 1.7 were estimated by using historical repository data; all faults detected for these versions were fixed for current versions of SpiderMonkey. We therefore performed a search for an identifying r for each test case t : r is the revision number of the first commit to the source repository such that (1) t fails for the version of the code before commit r

1. In C/C++, $a[i]$ is syntactic sugar for $*(a+i)$.




Fig. 1. SpiderMonkey 1.6 Results

and (2) t no longer fails once commit r is included. We assume commit r fixes t and therefore identifies the underlying fault exposed by t . This method is obviously approximate. Multiple faults may be fixed in one commit, or something that we would conceptually call “one fault” may be fixed by multiple partial fixes that do not completely handle the problem but handle some failing tests (e.g., a comparison operator such as > 10 that should be instead > 8 is changed to > 9). Additionally, a non-fix may sometimes cause a test to behave differently — for example, introducing a new optimization may sometimes cause tests exposing a still-present fault to no longer trigger it, as the fault-inducing aspect of the compiled code is removed by the new optimization. However, hand examination of a similar method in previous work [8] showed that in general this is a very good approximation of actual faults. For SpiderMonkey 1.8.5, too few of the discovered faults were clearly fixed in the latest releases, the source code repository mechanism changed, and the number of distinct faults appeared to be small enough to make the uncertainties of this method problematic, so we only measured actual failures, as these were themselves quite rare for 1.8.5.

1.2 Results and Statistical Validation

In order to statistically validate these results, we applied two-tailed Welch’s t -test and Wilcoxon U -test to the 96 individual 30-minute test blocks for swarm and default. For all measures (branch coverage, statement coverage,

TABLE 1
95% effect sizes (absolute) for swarm over default, SpiderMonkey

Version	ST	BR	Failures	Faults
1.6	40.4 - 72.0	25.8 - 54.0	16.3 - 22.4	2.6 - 4.9
1.7	-7.5 - 37.6	27.2 - 51.1	15.9 - 19.1	3.7 - 4.9
1.8.5	63.3 - 90.4	44.6 - 70.1	0.3 - 1.2	N/A

failures, and faults) and all three SpiderMonkey versions, the differences between swarm and the default strategy were statistically significant at the 99% level (in fact, the largest p -value observed was 0.003 for 1.7 statements). Table 1 shows 95% confidence interval bounds from the Welch test on the effect sizes for each measure, for all versions. Note that while the absolute difference for failures for 1.8.5 is small, only 1.6 failures on average were found using the default strategy during each 30 minute test run, so 0.3 failures is a nearly 19% improvement. The absolute coverage differences represent small relative improvements, while the small lower bounds on absolute distinct fault improvements for versions 1.6 and 1.7 are *nearly 40% and nearly 70% improvements*, respectively over the default. While the coverage improvements are also desirable, a 40-70% improvement in fault detection for a trivial-to-implement modification to test generation is the key argument for swarm testing’s value.

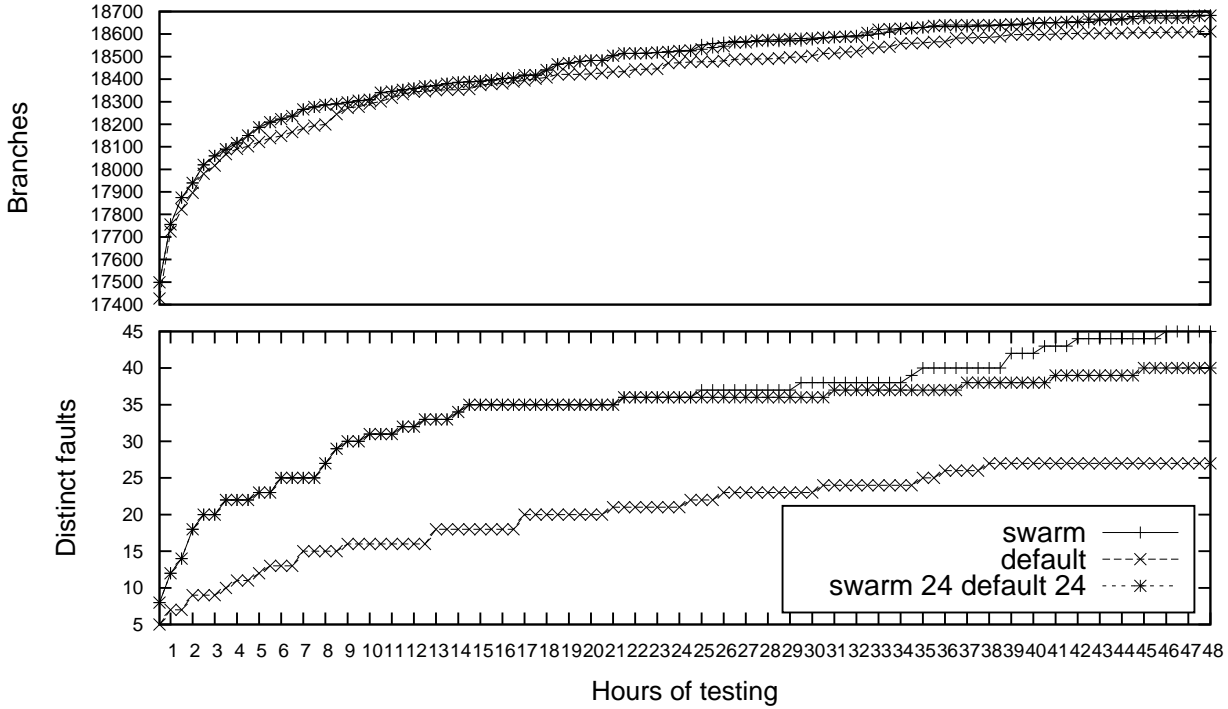


Fig. 2. SpiderMonkey 1.7 Results

1.3 Using Feedback to Improve Swarm

Figures 4 and 5 show the results of attempting to exploit the results of the first 24 hours of swarm testing to improve testing for the next 24 hour period, tuning the configurations used.

Because measuring code coverage after each test is expensive, we used a feedback approach based on coverage over difficult-to-execute coverage entities only. We considered a branch or line “uninteresting” if it was executed during every 30 minute run in the first 24 hours of testing. Coverage for the remaining “interesting” blocks and branches only was collected for each test, and ranked configurations by the total value (the log of the inverse of the frequency of coverage over all tests) of all covered entities. The feedback strategies were then (1) to re-use the best (or best 50, 100, 200, or 400) configurations in round-robin fashion for future test generation, and (2) to include each feature with a biased non-50% probability based on its frequency in tests executing at least one interesting coverage entity.

While some strategies show a cumulative improvement on swarm in terms of coverage or fault detection, *no improvements* (measured in terms of 30 minute run improvements on the cumulative swarm coverage at 24 hours) are statistically significant; in some cases there was a statistically significant *reduction* in new distinct faults per run. The results were similar enough to the pure swarm strategy, and required sufficient computational effort (in computing detailed coverage results and ranking configurations) and additional test infrastructure development that we find little support for applying simple feedback mechanisms. It is possible more sophis-

ticated methods, e.g. using machine learning, may be valuable, but the most important goal would be to cover *never-executed* code, for which there is, in general, no data from which learn a good configuration. Our hypothesis was that executing infrequently executed code more often should lead to either detection of bugs concerning that code, or execution of “nearby” uncovered code. While feedback did indeed increase the frequency of rarely executed entities, it did not produce a clear effect in terms of bugs or novel coverage.

2 YAFFS2 RESULTS

YAFFS2 [9] is a popular open-source NAND flash file system for embedded use; it was the default image format for early versions of the Android operating system. The test generator for YAFFS uses 47 different API calls as features. Applying swarm testing to YAFFS2 was essentially trivial, as the capability to turn on or off API calls is natural to such API-based random test generators; we did refactor the control from a `#define` based approach to command-line parameters to ease use of swarm testing.

Figure 6 shows that swarm testing improves not only branch and statement coverage but also mutation kill rates for YAFFS2. Because YAFFS2 is a smaller and simpler program than SpiderMonkey, we used 10 minute rather than 30 minute test intervals. All differences between swarm and default testing were statistically significant by Wilcoxon U-test for the 10 minute intervals, with p -values of 3.1×10^{-7} or lower. Tables 2 summarizes 95% confidence interval absolute gains in coverage and

TABLE 2
95% effect sizes (absolute) for swarm over default,
YAFFS2

ST	BR	Mutants Killed
15.7 - 50.1	10.5 - 28.1	

mutation kill for 10 minutes of swarm testing vs. 10 minutes of default testing.

3 CONCLUSION

Swarm testing relies on the following claim: for realistic systems, *randomly excluding some features from some tests* can improve coverage and fault detection, compared to a test suite that potentially uses every feature in every test. The benefit of using of a single, inclusive, default configuration—that every test can potentially expose any fault and cover any behavior, heretofore usually taken for granted in random testing—does not, in practice, make up for the fact that *some features can, statistically, suppress behaviors*. Effective testing therefore may require feature omission diversity. We show that this not only holds for simple container-class examples (e.g., pop operations suppress stack overflow) but for a widely used flash file system and 14 out of 17 versions of five production-quality C compilers. For these real-world systems, if we compare testing with a single inclusive configuration to testing with a set of 100–1,000 unique configurations, each omitting features with 50% probability per feature, we have observed (1) significantly better fault detection, (2) significantly better branch and statement coverage, and (3) strictly superior mutant detection. Test configuration diversity does indeed produce better testing in many realistic situations.

Swarm testing was inspired by swarm verification, and we hope that its ideas can be ported back to model checking. We also plan to investigate swarm in the context of bounded exhaustive testing and learning-based testing methods. Finally, we believe there is room to better understand *why* swarm provides its benefits, particularly in the context of large, idiosyncratic SUTs such as compilers, virtual machines, and OS kernels. More case studies will be needed to generate data to support this work. We also plan to investigate how swarm testing’s increased diversity of code coverage in test cases can benefit fault localization and program understanding algorithms relying on test cases [10]; traditional random tests are far more homogeneous than swarm tests.

We have made Python scripts supporting swarm testing available at <http://beaversource.oregonstate.edu/projects/cswarm/browser/release>.

REFERENCES

- [1] R. Hamlet, “Random testing,” in *Encyclopedia of Software Engineering*. Wiley, 1994, pp. 970–978.
- [2] J. H. Andrews, F. C. H. Li, and T. Menzies, “Nighthawk: A two-level genetic-random unit test data generator,” in *Proc. ASE*, Nov. 2007, pp. 144–153.
- [3] J. H. Andrews, A. Groce, M. Weston, and R.-G. Xu, “Random test run length and effectiveness,” in *Proc. ASE*, Sep. 2008, pp. 19–28.
- [4] X. Yang, Y. Chen, E. Eide, and J. Regehr, “Finding and understanding bugs in C compilers,” in *Proc. PLDI*, Jun. 2011, pp. 283–294.
- [5] A. Groce, G. Holzmann, and R. Joshi, “Randomized differential testing as a prelude to formal verification,” in *Proc. ICSE*, May 2007, pp. 621–631.
- [6] C. Pacheco, S. K. Lahiri, M. D. Ernst, and T. Ball, “Feedback-directed random test generation,” in *Proc. ICSE*, May 2007, pp. 75–84.
- [7] J. Ruderman, “Introducing jsfunfuzz,” 2007, <http://www.squarefree.com/2007/08/02/introducing-jsfunfuzz/>.
- [8] Y. Chen, A. Groce, C. Zhang, W.-K. Wong, X. Fern, E. Eide, and J. Regehr, “Taming compiler fuzzers,” in *Proc. PLDI*, Jun. 2013, pp. 197–208.
- [9] “YAFFS: A flash file system for embedded use,” <http://www.yaffs.net/>.
- [10] J. A. Jones and M. J. Harrold, “Empirical evaluation of the Tarantula automatic fault-localization technique,” in *Proc. ASE*, Nov. 2005, pp. 273–282.




Fig. 3. SpiderMonkey 1.8.5 Results: Note that this graph shows failures rather than faults.




Fig. 4. SpiderMonkey 1.6 Feedback Results




Fig. 5. SpiderMonkey 1.7 Feedback Results




Fig. 6. YAFFS2 Results